

Normalization Techniques

Devansh Arpit

Table of Contents

- 1 Introduction
- 2 Motivation
- 3 Batch Normalization
- 4 Normalization Propagation
- 5 Weight Normalization
- 6 Layer Normalization
- 7 Conclusion

- Task:

Train deep networks using gradient descent based methods

- Task:

Train deep networks using gradient descent based methods

- Goal:


Improve optimization for faster convergence

Table of Contents


- 1 Introduction
- 2 Motivation**
- 3 Batch Normalization
- 4 Normalization Propagation
- 5 Weight Normalization
- 6 Layer Normalization
- 7 Conclusion

- Transfer mapping from domain 1 to domain 2:

$$\mathbf{X} \xrightarrow{P(\mathbf{Y}|\mathbf{X})} \mathbf{Y}$$



Domain 1




Domain 2


- $P(\mathbf{Y}|\mathbf{X})$ identical for both Domain
- Learn best approximation $P_{\theta}(\mathbf{Y}|\mathbf{X}) \approx P(\mathbf{Y}|\mathbf{X})$

- Transfer mapping from domain 1 to domain 2:

$$\mathbf{X} \xrightarrow{P(\mathbf{Y}|\mathbf{X})} \mathbf{Y}$$



Domain 1




Domain 2


- $P(\mathbf{Y}|\mathbf{X})$ identical for both Domain
- Learn best approximation $P_{\theta}(\mathbf{Y}|\mathbf{X}) \approx P(\mathbf{Y}|\mathbf{X})$
- Maximum Likelihood focuses on high density regions of \mathbf{X}

- Transfer mapping from domain 1 to domain 2:

$$\mathbf{X} \xrightarrow{P(\mathbf{Y}|\mathbf{X})} \mathbf{Y}$$



Domain 1




Domain 2


- $P(\mathbf{Y}|\mathbf{X})$ identical for both Domain
- Learn best approximation $P_{\theta}(\mathbf{Y}|\mathbf{X}) \approx P(\mathbf{Y}|\mathbf{X})$
- Maximum Likelihood focuses on high density regions of \mathbf{X}
- Low density regions of \mathbf{X} are artifacts in $P_{\theta^*}(\mathbf{Y}|\mathbf{X})$

- Transfer mapping from domain 1 to domain 2:

$$\mathbf{X} \xrightarrow{P(\mathbf{Y}|\mathbf{X})} \mathbf{Y}$$



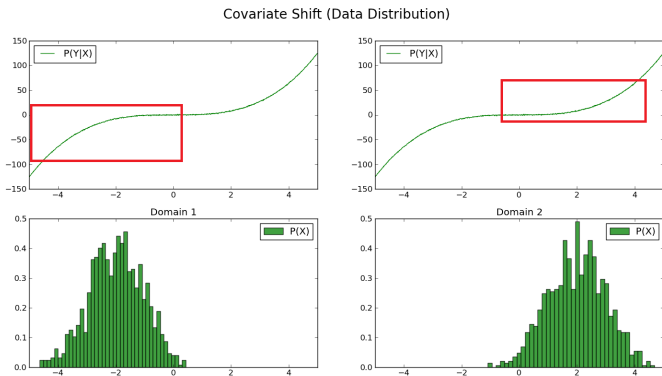
Domain 1



Domain 2

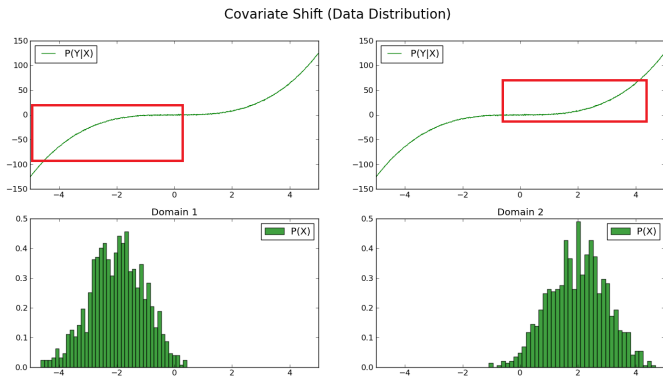
- $P(\mathbf{Y}|\mathbf{X})$ identical for both Domain
- Learn best approximation $P_{\theta}(\mathbf{Y}|\mathbf{X}) \approx P(\mathbf{Y}|\mathbf{X})$
- Maximum Likelihood focuses on high density regions of \mathbf{X}
- Low density regions of \mathbf{X} are artifacts in $P_{\theta^*}(\mathbf{Y}|\mathbf{X})$
- $P_{\theta^*}(\mathbf{Y}|\mathbf{X})$ for Domain 1 $\neq P_{\theta^*}(\mathbf{Y}|\mathbf{X})$ for Domain 2

- Simulation:



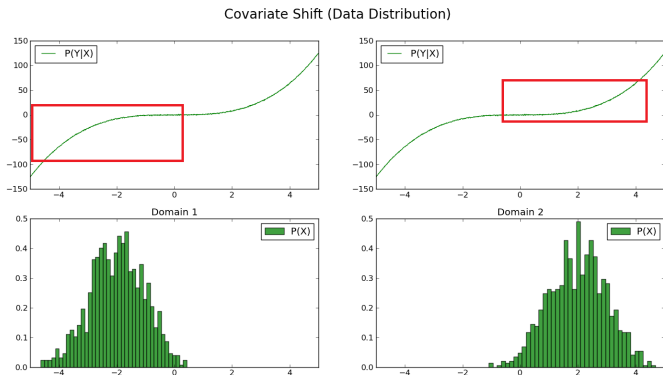
- Design: $P(X)$ is different but $P(Y|X)$ is identical

- Simulation:



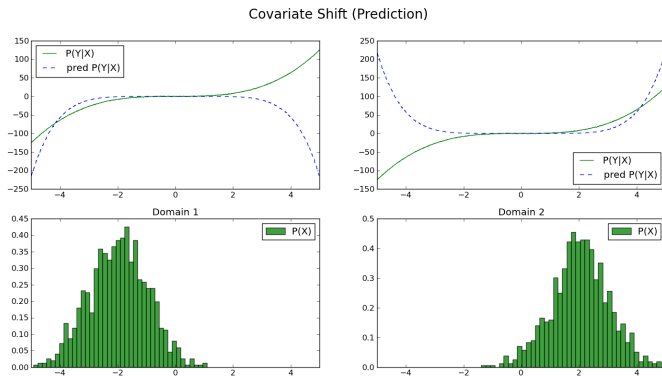
- Design: $P(X)$ is different but $P(Y|X)$ is identical
- Sample $x \sim P(X)$ separately from the 2 domains, then generate y for each x using $P(Y|X)$

- Simulation:




- Design: $P(X)$ is different but $P(Y|X)$ is identical
- Sample $x \sim P(X)$ separately from the 2 domains, then generate y for each x using $P(Y|X)$
- Sampled points for each domain are along the curve contained in the red box

- Simulation:



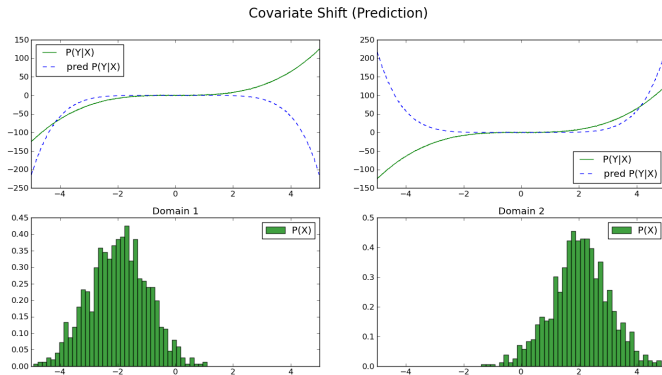
- Learn $P_{\theta}(Y|X)$ separately using the sampled (x,y) pair

- Simulation:



- Learn $P_{\theta}(Y|X)$ separately using the sampled (x,y) pair
- Learned $P_{\theta^*}(Y|X)$ for both domains are shown by dotted curves

- Simulation:




- Learn $P_{\theta}(Y|X)$ separately using the sampled (x,y) pair
- Learned $P_{\theta^*}(Y|X)$ for both domains are shown by dotted curves
- Optimal $P_{\theta^*}(Y|X)$ are different $\implies P_{\theta^*}(Y|X)$ not transferable

Motivation 1/2– Covariate Shift

- Internal Covariate Shift (Ioffe and Szegedy, 2015)

$$\mathbf{X}_i \xrightarrow{P(\mathbf{X}_{i+1}|\mathbf{X}_i)} \mathbf{X}_{i+1}$$

- Multi-layer end-to-end learning model




- Learn the best approximation $P_{\theta}(\mathbf{X}_{i+1}|\mathbf{X}_i) \approx P(\mathbf{X}_{i+1}|\mathbf{X}_i)$

Motivation 1/2– Covariate Shift

- Internal Covariate Shift (Ioffe and Szegedy, 2015)

$$\mathbf{X}_i \xrightarrow{P(\mathbf{X}_{i+1}|\mathbf{X}_i)} \mathbf{X}_{i+1}$$

- Multi-layer end-to-end learning model




- Learn the best approximation $P_\theta(\mathbf{X}_{i+1}|\mathbf{X}_i) \approx P(\mathbf{X}_{i+1}|\mathbf{X}_i)$
- During SGD updates, hidden layer $P(\mathbf{X}_i)$ keeps shifting

Motivation 1/2– Covariate Shift

- Internal Covariate Shift (Ioffe and Szegedy, 2015)

$$\mathbf{X}_i \xrightarrow{P(\mathbf{X}_{i+1}|\mathbf{X}_i)} \mathbf{X}_{i+1}$$

- Multi-layer end-to-end learning model




- Learn the best approximation $P_{\theta}(\mathbf{X}_{i+1}|\mathbf{X}_i) \approx P(\mathbf{X}_{i+1}|\mathbf{X}_i)$
- During SGD updates, hidden layer $P(\mathbf{X}_i)$ keeps shifting
- \implies target $P_{\theta^*}(\mathbf{X}_{i+1}|\mathbf{X}_i)$ keeps shifting

Motivation 1/2– Covariate Shift

- Internal Covariate Shift (Ioffe and Szegedy, 2015)

$$\mathbf{X}_i \xrightarrow{P(\mathbf{X}_{i+1}|\mathbf{X}_i)} \mathbf{X}_{i+1}$$

- Multi-layer end-to-end learning model




- Learn the best approximation $P_\theta(\mathbf{X}_{i+1}|\mathbf{X}_i) \approx P(\mathbf{X}_{i+1}|\mathbf{X}_i)$
- During SGD updates, hidden layer $P(\mathbf{X}_i)$ keeps shifting
- \implies target $P_{\theta^*}(\mathbf{X}_{i+1}|\mathbf{X}_i)$ keeps shifting
- Learning $P_\theta(\mathbf{X}_{i+1}|\mathbf{X}_i)$ using SGD is slow

What happens if input samples are not 0 mean?

- Let: $h_i = \sigma(a_i)$, where
 $a_i = \mathbf{w}_i^T \mathbf{x} + b_i$
- Consider the SGD weight update equation:


$$\begin{aligned}\mathbf{w}_i^{t+1} &= \mathbf{w}_i^t - \eta \frac{\partial \mathcal{L}}{\partial \mathbf{w}_i} \\ &= \mathbf{w}_i^t - \eta \frac{\partial \mathcal{L}}{\partial a_i} \frac{\partial a_i}{\partial \mathbf{w}_i} \\ &= \mathbf{w}_i^t - \eta \delta \mathbf{x}\end{aligned}$$



What happens if input samples are not 0 mean?

- Let: $h_i = \sigma(a_i)$, where
 $a_i = \mathbf{w}_i^T \mathbf{x} + b_i$
- Consider the SGD weight update equation:

$$\begin{aligned}\mathbf{w}_i^{t+1} &= \mathbf{w}_i^t - \eta \frac{\partial \mathcal{L}}{\partial \mathbf{w}_i} \\ &= \mathbf{w}_i^t - \eta \frac{\partial \mathcal{L}}{\partial a_i} \frac{\partial a_i}{\partial \mathbf{w}_i} \\ &= \mathbf{w}_i^t - \eta \delta \mathbf{x}\end{aligned}$$




- If all \mathbf{x} are positive \implies all the weight units get updated with the same sign!

What happens if input samples are not 0 mean?

- Let: $h_i = \sigma(a_i)$, where
 $a_i = \mathbf{w}_i^T \mathbf{x} + b_i$
- Consider the SGD weight update equation:

$$\begin{aligned}\mathbf{w}_i^{t+1} &= \mathbf{w}_i^t - \eta \frac{\partial \mathcal{L}}{\partial \mathbf{w}_i} \\ &= \mathbf{w}_i^t - \eta \frac{\partial \mathcal{L}}{\partial a_i} \frac{\partial a_i}{\partial \mathbf{w}_i} \\ &= \mathbf{w}_i^t - \eta \delta \mathbf{x}\end{aligned}$$




- If all \mathbf{x} are positive \implies all the weight units get updated with the same sign!
- Any bias in \mathbf{x} introduces bias in weight updates– bad!
- Learning slows down– path to optimal \mathbf{w}^* becomes longer

What happens if input samples are not 0 mean and unit variance?
(A more concrete special case)

- Consider loss:

$$\mathcal{L}(\mathbf{W}) = (1/2)\mathbb{E}_{\mathbf{x}, \mathbf{y}}[\|\mathbf{y} - \mathbf{W}\mathbf{x}\|^2]$$

- Update rule: $\mathbf{W}_i^{t+1} = \mathbf{W}_i^t - \eta \frac{\partial \mathcal{L}}{\partial \mathbf{W}_i}$




What happens if input samples are not 0 mean and unit variance?
(A more concrete special case)

- Consider loss:

$$\mathcal{L}(\mathbf{W}) = (1/2)\mathbb{E}_{\mathbf{x},\mathbf{y}}[\|\mathbf{y} - \mathbf{W}\mathbf{x}\|^2]$$

- Update rule: $\mathbf{W}_i^{t+1} = \mathbf{W}_i^t - \eta \frac{\partial \mathcal{L}}{\partial \mathbf{W}_i}$



- Using Taylor's expansion around a minima \mathbf{W}_i^* :


$$\mathcal{L}(\mathbf{W}_i) = \mathcal{L}(\mathbf{W}_i^*) + (\mathbf{W}_i - \mathbf{W}_i^*)^T \frac{\partial \mathcal{L}}{\partial \mathbf{W}_i^*} (\mathbf{W}_i - \mathbf{W}_i^*)$$

What happens if input samples are not 0 mean and unit variance?
(A more concrete special case)

- Consider loss:

$$\mathcal{L}(\mathbf{W}) = (1/2)\mathbb{E}_{\mathbf{x},\mathbf{y}}[\|\mathbf{y} - \mathbf{W}\mathbf{x}\|^2]$$

- Update rule: $\mathbf{W}_i^{t+1} = \mathbf{W}_i^t - \eta \frac{\partial \mathcal{L}}{\partial \mathbf{W}_i}$



- Using Taylor's expansion around a minima \mathbf{W}_i^* :

$$\mathcal{L}(\mathbf{W}_i) = \mathcal{L}(\mathbf{W}_i^*) + (\mathbf{W}_i - \mathbf{W}_i^*)^T \frac{\partial^2 \mathcal{L}}{\partial \mathbf{W}_i^2} (\mathbf{W}_i - \mathbf{W}_i^*)$$

- Taking derivative on both sides:


$$\frac{\partial \mathcal{L}(\mathbf{W}_i)}{\partial \mathbf{W}_i} = (1/2) \frac{\partial^2 \mathcal{L}}{\partial \mathbf{W}_i^2} (\mathbf{W}_i - \mathbf{W}_i^*) \quad \left(\frac{\partial^2 \mathcal{L}}{\partial \mathbf{W}_i^2} = \mathbb{E}_{\mathbf{x}}[\mathbf{x}\mathbf{x}^T] \right)$$

What happens if input samples are not 0 mean and unit variance?
(A more concrete special case)

- Consider loss:

$$\mathcal{L}(\mathbf{W}) = (1/2)\mathbb{E}_{\mathbf{x},\mathbf{y}}[\|\mathbf{y} - \mathbf{W}\mathbf{x}\|^2]$$

- Update rule: $\mathbf{W}_i^{t+1} = \mathbf{W}_i^t - \eta \frac{\partial \mathcal{L}}{\partial \mathbf{W}_i}$



- Using Taylor's expansion around a minima \mathbf{W}_i^* :

$$\mathcal{L}(\mathbf{W}_i) = \mathcal{L}(\mathbf{W}_i^*) + (\mathbf{W}_i - \mathbf{W}_i^*)^T \frac{\partial^2 \mathcal{L}}{\partial \mathbf{W}_i^2} (\mathbf{W}_i - \mathbf{W}_i^*)$$

- Taking derivative on both sides:

$$\frac{\partial \mathcal{L}(\mathbf{W}_i)}{\partial \mathbf{W}_i} = (1/2) \frac{\partial^2 \mathcal{L}}{\partial \mathbf{W}_i^2} (\mathbf{W}_i - \mathbf{W}_i^*) \quad \left(\frac{\partial^2 \mathcal{L}}{\partial \mathbf{W}_i^2} = \mathbb{E}_{\mathbf{x}}[\mathbf{x}\mathbf{x}^T] \right)$$

- Thus update rule is:

$$\mathbf{W}_i^{t+1} = \mathbf{W}_i^t - \eta \mathbb{E}_{\mathbf{x}}[\mathbf{x}\mathbf{x}^T] (\mathbf{W}_i^t - \mathbf{W}_i^*)$$




What happens if input samples are not 0 mean and unit variance?
(A more concrete special case)

- Consider loss:

$$\mathcal{L}(\mathbf{W}) = (1/2)\mathbb{E}_{\mathbf{x},\mathbf{y}}[\|\mathbf{y} - \mathbf{W}\mathbf{x}\|^2]$$



- Update rule:

$$\mathbf{W}_i^{t+1} = \mathbf{W}_i^t - \eta \mathbb{E}_{\mathbf{x}}[\mathbf{x}\mathbf{x}^T](\mathbf{W}_i^t - \mathbf{W}_i^*)$$



$$\mathbf{W}_i^{t+1} = \mathbf{W}_i^t - \eta \begin{pmatrix} \lambda_{\max} & & \\ & \ddots & \\ & & \lambda_{\min} \end{pmatrix} (\mathbf{W}_i^t - \mathbf{W}_i^*)$$

What happens if input samples are not 0 mean and unit variance?
(A more concrete special case)



$$\mathbf{W}_i^{t+1} = \mathbf{W}_i^t - \eta \begin{pmatrix} \lambda_{\max} & & \\ & \ddots & \\ & & \lambda_{\min} \end{pmatrix} (\mathbf{W}_i^t - \mathbf{W}_i^*)$$

- Large mean and unequal variance $\implies \lambda_{\max} \gg \lambda_{\min}$
- Each dimension of \mathbf{W}_i^t needs a different learning rate
- Hence overall learning rate η is capped by $\frac{1}{\lambda_{\max}}$

(λ_{\max} is the largest eigenvalue of $\mathbb{E}_{\mathbf{x}}[\mathbf{x}\mathbf{x}^T]$)

Take home message:

- At least have zero mean (avoid single large eigenvalue)
- Equal variance across dimensions is better (not very useful if mean $\neq 0$)
- Uncorrelated (spherical) representations are best (but expensive!)

Take home message:

- At least have zero mean (avoid single large eigenvalue)
- Equal variance across dimensions is better (not very useful if mean $\neq 0$)
- Uncorrelated (spherical) representations are best (but expensive!)
- Maintain these properties at all hidden layers during training!

Table of Contents

- 1 Introduction
- 2 Motivation
- 3 Batch Normalization**
- 4 Normalization Propagation
- 5 Weight Normalization
- 6 Layer Normalization
- 7 Conclusion

- Avoid covariate shift
 - Normalize the entire distribution at every layer at every iteration!

- Avoid covariate shift
 - Normalize the entire distribution at every layer at every iteration!
 - Simply infeasible!!!

- Avoid covariate shift
 - Normalize the entire distribution at every layer at every iteration!
 - Simply infeasible!!!
- Alleviate it by normalizing the 1st two orders of statistics?
 - Needs whitening over the entire dataset at every iteration!

- Avoid covariate shift
 - Normalize the entire distribution at every layer at every iteration!
 - Simply infeasible!!!
- Alleviate it by normalizing the 1st two orders of statistics?
 - Needs whitening over the entire dataset at every iteration!
 - Still very expensive!!!

- Avoid covariate shift
 - Normalize the entire distribution at every layer at every iteration!
 - Simply infeasible!!!
- Alleviate it by normalizing the 1st two orders of statistics?
 - Needs whitening over the entire dataset at every iteration!
 - Still very expensive!!!
- Key idea:
 - Compute statistics over mini-batch
 - Compute statistics unit-wise (notice mean is still optimal per batch)


- Avoid covariate shift
 - Normalize the entire distribution at every layer at every iteration!
 - Simply infeasible!!!
- Alleviate it by normalizing the 1st two orders of statistics?
 - Needs whitening over the entire dataset at every iteration!
 - Still very expensive!!!
- Key idea:
 - Compute statistics over mini-batch
 - Compute statistics unit-wise (notice mean is still optimal per batch)Finally we can do something about it :)

- Consider any hidden unit's pre-activation a_i
- Then normalized pre-activation under BN is given by:

$$\text{BN}(a_i) = \frac{(\mathbf{a}_i - \mathbb{E}_{\mathcal{B}}[\mathbf{a}_i]))}{\sqrt{\text{var}_{\mathcal{B}}(\mathbf{a}_i)}}$$

- Consider any hidden unit's pre-activation a_i
- Then normalized pre-activation under BN is given by:

$$\text{BN}(a_i) = \frac{(\mathbf{a}_i - \mathbb{E}_{\mathcal{B}}[\mathbf{a}_i])}{\sqrt{\text{var}_{\mathcal{B}}(\mathbf{a}_i)}}$$




- Pretty close!

(approximation is worst when principal components are maximally away from axis)

- A traditional vs. Batch Normalized (BN) ReLU layer:

Traditional:




BN:

$$\mathbf{x} \quad \hat{a}_i = \frac{\gamma_i (\mathbf{W}_i^T (\mathbf{x} - \mathbb{E}_{\mathcal{B}}[\mathbf{x}]))}{\sqrt{\text{var}_{\mathcal{B}}(\mathbf{W}_i^T \mathbf{x})}} + \beta_i \quad \mathbf{h} = \text{ReLU}(\hat{\mathbf{a}}) \quad \mathbf{o} = \mathbf{h}$$

- A traditional vs. Batch Normalized (BN) ReLU layer:

Traditional:




BN:

$$\mathbf{x} \quad \hat{a}_i = \frac{\gamma_i (\mathbf{W}_i^T (\mathbf{x} - \mathbb{E}_{\mathcal{B}}[\mathbf{x}]))}{\sqrt{\text{var}_{\mathcal{B}}(\mathbf{W}_i^T \mathbf{x})}} + \beta_i \quad \mathbf{h} = \text{ReLU}(\hat{\mathbf{a}}) \quad \mathbf{o} = \mathbf{h}$$

- γ_i and β_i are learnable scale and bias parameters

- A traditional vs. Batch Normalized (BN) ReLU layer:

Traditional:




BN:

$$\mathbf{x} \quad \hat{a}_i = \frac{\gamma_i (\mathbf{W}_i^T (\mathbf{x} - \mathbb{E}_{\mathcal{B}}[\mathbf{x}]))}{\sqrt{\text{var}_{\mathcal{B}}(\mathbf{W}_i^T \mathbf{x})}} + \beta_i \quad \mathbf{h} = \text{ReLU}(\hat{\mathbf{a}}) \quad \mathbf{o} = \mathbf{h}$$

- γ_i and β_i are learnable scale and bias parameters
- Notice bias β_i is outside the normalization (why?)

- A traditional vs. Batch Normalized (BN) ReLU layer:

Traditional:




BN:

$$\mathbf{x} \quad \hat{a}_i = \frac{\gamma_i (\mathbf{W}_i^T (\mathbf{x} - \mathbb{E}_{\mathcal{B}}[\mathbf{x}]))}{\sqrt{\text{var}_{\mathcal{B}}(\mathbf{W}_i^T \mathbf{x})}} + \beta_i \quad \mathbf{h} = \text{ReLU}(\hat{\mathbf{a}}) \quad \mathbf{o} = \mathbf{h}$$

- γ_i and β_i are learnable scale and bias parameters
- Notice bias β_i is outside the normalization (why?)
- Why not apply BN post-activation ? (hint: Gaussianity)

- A traditional vs. Batch Normalized (BN) ReLU layer:

Traditional:




BN:


$$\mathbf{x} \quad \hat{\mathbf{a}}_i = \frac{\gamma_i (\mathbf{W}_i^T (\mathbf{x} - \mathbb{E}_{\mathcal{B}}[\mathbf{x}]))}{\sqrt{\text{var}_{\mathcal{B}}(\mathbf{W}_i^T \mathbf{x})}} + \beta_i \quad \mathbf{h} = \text{ReLU}(\hat{\mathbf{a}}) \quad \mathbf{o} = \mathbf{h}$$

- γ_i and β_i are learnable scale and bias parameters
- Notice bias β_i is outside the normalization (why?)
- Why not apply BN post-activation ? (hint: Gaussianity)
- During test time— use running average estimates of mean and standard deviation

- How to extend to convolutional layers?
Simply treat each depth vector as a separate sample



- How to extend to convolutional layers?
Simply treat each depth vector as a separate sample



- Back-propagate through the normalization
Otherwise normalizing representations changes prediction

Effects of batch normalization:

- Parameter scaling:
 - Representations become scale invariant
 $\text{BN}(c\mathbf{W}\mathbf{x}) = \text{BN}(\mathbf{W}\mathbf{x})$
 - Gradients become inversely proportional to parameter scale
 $\frac{\partial \text{BN}(c\mathbf{W}\mathbf{x})}{\partial c\mathbf{W}} = (1/c) \frac{\partial \text{BN}(\mathbf{W}\mathbf{x})}{\partial \mathbf{W}}$
Allows for large learning rates!

Effects of batch normalization:

- Parameter scaling:
 - Representations become scale invariant
$$\text{BN}(c\mathbf{W}\mathbf{x}) = \text{BN}(\mathbf{W}\mathbf{x})$$
 - Gradients become inversely proportional to parameter scale
$$\frac{\partial \text{BN}(c\mathbf{W}\mathbf{x})}{\partial c\mathbf{W}} = (1/c) \frac{\partial \text{BN}(\mathbf{W}\mathbf{x})}{\partial \mathbf{W}}$$
Allows for large learning rates!
- Regularization:
 - Each sample gets a different representation depending on mini-batch

Effects of batch normalization:

- Parameter scaling:
 - Representations become scale invariant
$$\text{BN}(c\mathbf{W}\mathbf{x}) = \text{BN}(\mathbf{W}\mathbf{x})$$
 - Gradients become inversely proportional to parameter scale
$$\frac{\partial \text{BN}(c\mathbf{W}\mathbf{x})}{\partial c\mathbf{W}} = (1/c) \frac{\partial \text{BN}(\mathbf{W}\mathbf{x})}{\partial \mathbf{W}}$$
Allows for large learning rates!
- Regularization:
 - Each sample gets a different representation depending on mini-batch

Awesome trick!

But also puzzling at the same time—

What if 2 different samples in two different batches get the same representation after BN?

Batch Normalization (Ioffe and Szegedy, 2015)




Figure 2: Single crop validation accuracy of Inception and its batch-normalized variants, vs. the number of training steps.

Model	Steps to 72.2%	Max accuracy
Inception	$31.0 \cdot 10^6$	72.2%
BN-Baseline	$13.3 \cdot 10^6$	72.7%
BN-x5	$2.1 \cdot 10^6$	73.0%
BN-x30	$2.7 \cdot 10^6$	74.8%
BN-x5-Sigmoid		69.8%

Figure 3: For Inception and the batch-normalized variants, the number of training steps required to reach the maximum accuracy of Inception (72.2%), and the maximum accuracy achieved by the network.

Figures taken from Ioffe, Sergey, and Christian Szegedy. "Batch normalization: Accelerating deep network training by reducing internal covariate shift." ICML (2015).

Table of Contents

- 1 Introduction
- 2 Motivation
- 3 Batch Normalization
- 4 Normalization Propagation**
- 5 Weight Normalization
- 6 Layer Normalization
- 7 Conclusion

Key idea:

- Instead of recomputing statistics at every layer, exploit normalization in data by propagating it to hidden layers

$$\text{Coherence} = \max_{\mathbf{w}_i, \mathbf{w}_j, i \neq j} \frac{|\mathbf{w}_i^T \mathbf{w}_j|}{\|\mathbf{w}_i\|_2 \|\mathbf{w}_j\|_2}$$

Incoherence \implies pairwise angle between vectors is large

Key idea:


- Instead of recomputing statistics at every layer, exploit normalization in data by propagating it to hidden layers
- Trick: If \mathbf{x} has $\mathbf{0}$ mean \mathbf{I} covariance, then $\mathbf{W}\mathbf{x}$ also has $\mathbf{0}$ mean $\approx \mathbf{I}$ covariance if $\|\mathbf{W}_i\|_2 = 1$ and \mathbf{W} is incoherent

$$\text{Coherence} = \max_{\mathbf{w}_i, \mathbf{w}_j, i \neq j} \frac{|\mathbf{w}_i^T \mathbf{w}_j|}{\|\mathbf{w}_i\|_2 \|\mathbf{w}_j\|_2}$$

Incoherence \implies pairwise angle between vectors is large

Key idea:

- Instead of recomputing statistics at every layer, exploit normalization in data by propagating it to hidden layers
- Trick: If \mathbf{x} has $\mathbf{0}$ mean \mathbf{I} covariance, then $\mathbf{W}\mathbf{x}$ also has $\mathbf{0}$ mean $\approx \mathbf{I}$ covariance if $\|\mathbf{W}_i\|_2 = 1$ and \mathbf{W} is incoherent




$$\text{Coherence} = \max_{\mathbf{W}_i, \mathbf{W}_j, i \neq j} \frac{|\mathbf{W}_i^T \mathbf{W}_j|}{\|\mathbf{W}_i\|_2 \|\mathbf{W}_j\|_2}$$

Incoherence \implies pairwise angle between vectors is large

Key idea:

- Instead of recomputing statistics at every layer, exploit normalization in data by propagating it to hidden layers
- Trick: If \mathbf{x} has $\mathbf{0}$ mean \mathbf{I} covariance, then $\mathbf{W}\mathbf{x}$ also has $\mathbf{0}$ mean $\approx \mathbf{I}$ covariance if $\|\mathbf{W}_i\|_2 = 1$ and \mathbf{W} is incoherent




- Assumption: pre-activations $\mathbf{W}\mathbf{x}$ are approximately Gaussian $\sigma(\mathbf{W}_i^T \mathbf{x})$ has fixed mean and variance

$$\text{Coherence} = \max_{\mathbf{W}_i, \mathbf{W}_j, i \neq j} \frac{|\mathbf{W}_i^T \mathbf{W}_j|}{\|\mathbf{W}_i\|_2 \|\mathbf{W}_j\|_2}$$

Incoherence \implies pairwise angle between vectors is large

- A traditional vs. NormProp ReLU layer:

Traditional:




NormProp:

Diagram illustrating a NormProp ReLU layer structure:

Input \mathbf{x} is connected to hidden nodes $\hat{a}_i = \frac{\gamma_i (\mathbf{W}_i^T \mathbf{x})}{\|\mathbf{W}_i\|_2} + \beta_i$ (blue circles). The hidden nodes are connected to output nodes $\mathbf{h} = \text{ReLU}(\hat{\mathbf{a}})$ (blue circles). The final output is $\mathbf{o} = \frac{\left[\mathbf{h} - \sqrt{\frac{1}{2\pi}} \right]}{\sqrt{\frac{1}{2} \left(1 - \frac{1}{\pi} \right)}}$ (blue circles). The diagram shows a fully connected layer from input to hidden, followed by a ReLU activation, and then a normalization step.

- A traditional vs. NormProp ReLU layer:

Traditional:




NormProp:

$$\mathbf{x} \quad \hat{\mathbf{a}}_i = \frac{\gamma_i (\mathbf{W}_i^T \mathbf{x})}{\|\mathbf{W}_i\|_2} + \beta_i \quad \mathbf{h} = \text{ReLU}(\hat{\mathbf{a}}) \quad \mathbf{o} = \frac{\left[\mathbf{h} - \sqrt{\frac{1}{2\pi}} \right]}{\sqrt{\frac{1}{2} \left(1 - \frac{1}{\pi} \right)}}$$

- γ_i and β_i are learnable scale and bias parameters

- A traditional vs. NormProp ReLU layer:

Traditional:




NormProp:

$$\mathbf{x} \quad \hat{a}_i = \frac{\gamma_i (\mathbf{W}_i^T \mathbf{x})}{\|\mathbf{W}_i\|_2} + \beta_i \quad \mathbf{h} = \text{ReLU}(\hat{\mathbf{a}}) \quad \mathbf{o} = \frac{\left[\mathbf{h} - \sqrt{\frac{1}{2\pi}} \right]}{\sqrt{\frac{1}{2} \left(1 - \frac{1}{\pi} \right)}}$$

- γ_i and β_i are learnable scale and bias parameters
- Data can be normalized globally or batch-wise

- A traditional vs. NormProp ReLU layer:

Traditional:




NormProp:

$$\mathbf{x} \quad \hat{a}_i = \frac{\gamma_i (\mathbf{W}_i^T \mathbf{x})}{\|\mathbf{W}_i\|_2} + \beta_i \quad \mathbf{h} = \text{ReLU}(\hat{\mathbf{a}}) \quad \mathbf{o} = \frac{\left[\mathbf{h} - \sqrt{\frac{1}{2\pi}} \right]}{\sqrt{\frac{1}{2} \left(1 - \frac{1}{\pi} \right)}}$$

- γ_i and β_i are learnable scale and bias parameters
- Data can be normalized globally or batch-wise
- train and test normalizations identical (for global data normalization)

- A traditional vs. NormProp ReLU layer:

Traditional:




NormProp:

$$\mathbf{x} \quad \hat{a}_i = \frac{\gamma_i (\mathbf{W}_i^T \mathbf{x})}{\|\mathbf{W}_i\|_2} + \beta_i \quad \mathbf{h} = \text{ReLU}(\hat{\mathbf{a}}) \quad \mathbf{o} = \frac{\left[\mathbf{h} - \sqrt{\frac{1}{2\pi}} \right]}{\sqrt{\frac{1}{2} \left(1 - \frac{1}{\pi} \right)}}$$

- γ_i and β_i are learnable scale and bias parameters
- Data can be normalized globally or batch-wise
- train and test normalizations identical (for global data normalization)
- Applicable with batch-size 1


- How to extend to convolutional layers?

$$\hat{a}_i = \frac{\gamma_i(\mathbf{W}_i * \mathbf{x})}{\|\mathbf{W}_i\|_F} + \beta_i \quad (i^{th} \text{ featuremap})$$



- How to extend to convolutional layers?

$$\hat{a}_i = \frac{\gamma_i(\mathbf{w}_i * \mathbf{x})}{\|\mathbf{w}_i\|_F} + \beta_i \quad (i^{th} \text{ featuremap})$$



- Back-propagate through the normalization of weight scale
Otherwise normalizing representations changes prediction


Analysis:

- Singular values of $\mathbf{J} = \frac{\partial \mathbf{o}}{\partial \mathbf{x}} \approx 1$ prevents gradient problems
(Saxe et al. 2014)
- For ReLU layer:
 - $\mathbb{E}_{\mathbf{x}}[\mathbf{J}\mathbf{J}^T] \approx 1.47\mathbf{I} \implies$ Singular values of $\mathbf{J} \approx 1.2$

Extension to other Activation functions (σ):

- $\mathbf{o}_i = \frac{1}{c_1} \left[\sigma \left(\frac{\gamma_i(\mathbf{W}_i * \mathbf{x})}{\|\mathbf{W}_i\|_F} + \beta_i \right) - c_2 \right]$
- $c_1 = \sqrt{\text{var}(\sigma(Y))}$, $c_2 = \mathbb{E}[\sigma(Y)]$
- Y has Standard Normal distribution

Identical parameter scaling and regularization effects as BN



Methods	Test Error (%)
CIFAR-10 with data augmentation	
NormProp	7.47
Batch Normalization	7.25
NIN + ALP units	7.51
NIN	8.81
DSN	7.97
Maxout	9.38

Table of Contents

- 1 Introduction
- 2 Motivation
- 3 Batch Normalization
- 4 Normalization Propagation
- 5 Weight Normalization**
- 6 Layer Normalization
- 7 Conclusion

Key idea:

- Decouple the scale and direction of weights

Key idea:

- Decouple the scale and direction of weights
- Initialize parameters to have 0 mean unit variance pre-activations

Key idea:

- Decouple the scale and direction of weights
- Initialize parameters to have 0 mean unit variance pre-activations
- Back-propagate through the normalization

- Consider any hidden unit's pre-activation a_i
- Then weight normalized pre-activation is given by:


$$\text{WeighNorm}(a_i) = \frac{\gamma_i(\mathbf{W}_i^T \mathbf{x})}{\|\mathbf{W}_i\|_2} + \beta_i$$

$$\text{Initialize: } \gamma_i = \frac{1}{\sqrt{\text{var}_{\mathcal{B}}(\mathbf{W}_i^T \mathbf{x}) / \|\mathbf{W}_i\|_2}} \quad \beta_i = -\frac{\mathbb{E}_{\mathcal{B}}[\mathbf{W}_i^T \mathbf{x} / \|\mathbf{W}_i\|_2]}{\sqrt{\text{var}_{\mathcal{B}}(\mathbf{W}_i^T \mathbf{x}) / \|\mathbf{W}_i\|_2}}$$

- Equivalently pre-activations get normalized using 1 mini-batch initially
- Optimization doesn't have explicit mean and variance normalization

- A traditional vs. Weight Normalized ReLU layer:

Traditional:




WeightNorm:

$\mathbf{x} \quad \hat{a}_i = \frac{\gamma_i(\mathbf{W}_i^T \mathbf{x})}{\|\mathbf{W}_i\|_2} + \beta_i \quad \mathbf{h} = \text{ReLU}(\hat{\mathbf{a}}) \quad \mathbf{o} = \mathbf{h}$

- A traditional vs. Weight Normalized ReLU layer:

Traditional:




WeightNorm:

$\mathbf{x} \quad \hat{\mathbf{a}}_i = \frac{\gamma_i (\mathbf{W}_i^T \mathbf{x})}{\|\mathbf{W}_i\|_2} + \beta_i \quad \mathbf{h} = \text{ReLU}(\hat{\mathbf{a}}) \quad \mathbf{o} = \mathbf{h}$

- γ_i and β_i are learnable scale and bias parameters

- A traditional vs. Weight Normalized ReLU layer:

Traditional:




WeightNorm:

$\mathbf{x} \quad \hat{a}_i = \frac{\gamma_i (\mathbf{W}_i^T \mathbf{x})}{\|\mathbf{W}_i\|_2} + \beta_i \quad \mathbf{h} = \text{ReLU}(\hat{\mathbf{a}}) \quad \mathbf{o} = \mathbf{h}$

- γ_i and β_i are learnable scale and bias parameters
- train and test normalizations identical

- A traditional vs. Weight Normalized ReLU layer:

Traditional:



WeightNorm:

$\mathbf{x} \quad \hat{a}_i = \frac{\gamma_i (\mathbf{W}_i^T \mathbf{x})}{\|\mathbf{W}_i\|_2} + \beta_i \quad \mathbf{h} = \text{ReLU}(\hat{\mathbf{a}}) \quad \mathbf{o} = \mathbf{h}$

- γ_i and β_i are learnable scale and bias parameters
- train and test normalizations identical
- Applicable with batch-size 1

Effect of backpropagating through normalization:

- Let $\theta = \frac{\gamma(\mathbf{w})}{\|\mathbf{w}\|_2}$, then using SGD:

$$\begin{aligned}\Delta \mathbf{w} &= -\eta \frac{\partial \mathcal{L}}{\partial \mathbf{w}} \\ &= -\eta \frac{\gamma}{\|\mathbf{w}^t\|_2} \left(\mathbf{I} - \frac{\theta^t \theta^{tT}}{\|\theta^t\|^2} \right) \frac{\partial \mathcal{L}}{\partial \theta}\end{aligned}$$

- Thus, $\mathbf{w} \perp \Delta \mathbf{w}$

Effect of backpropagating through normalization:

- Let $\theta = \frac{\gamma(\mathbf{w})}{\|\mathbf{w}\|_2}$, then using SGD:

$$\begin{aligned}\Delta \mathbf{w} &= -\eta \frac{\partial \mathcal{L}}{\partial \mathbf{w}} \\ &= -\eta \frac{\gamma}{\|\mathbf{w}^t\|_2} \left(\mathbf{I} - \frac{\theta^t \theta^{tT}}{\|\theta^t\|^2} \right) \frac{\partial \mathcal{L}}{\partial \theta}\end{aligned}$$

- Thus, $\mathbf{w} \perp \Delta \mathbf{w}$
- Let $c = \|\Delta \mathbf{w}\| / \|\mathbf{w}\|$, notice: $\|\mathbf{w}^{t+1}\| = \sqrt{(1 + c^2)} \|\mathbf{w}^t\|$

Effect of backpropagating through normalization:

- Let $\theta = \frac{\gamma(\mathbf{w})}{\|\mathbf{w}\|_2}$, then using SGD:

$$\begin{aligned}\Delta \mathbf{w} &= -\eta \frac{\partial \mathcal{L}}{\partial \mathbf{w}} \\ &= -\eta \frac{\gamma}{\|\mathbf{w}^t\|_2} \left(\mathbf{I} - \frac{\theta^t \theta^{tT}}{\|\theta^t\|^2} \right) \frac{\partial \mathcal{L}}{\partial \theta}\end{aligned}$$

- Thus, $\mathbf{w} \perp \Delta \mathbf{w}$
- Let $c = \|\Delta \mathbf{w}\| / \|\mathbf{w}\|$, notice: $\|\mathbf{w}^{t+1}\| = \sqrt{(1 + c^2)} \|\mathbf{w}^t\|$
- Large $c \implies$ large $\|\mathbf{w}^{t+1}\| \implies$ small gradient $\Delta \mathbf{w}$

Effect of backpropagating through normalization:

- Let $\theta = \frac{\gamma(\mathbf{w})}{\|\mathbf{w}\|_2}$, then using SGD:

$$\begin{aligned}\Delta \mathbf{w} &= -\eta \frac{\partial \mathcal{L}}{\partial \mathbf{w}} \\ &= -\eta \frac{\gamma}{\|\mathbf{w}^t\|_2} \left(\mathbf{I} - \frac{\theta^t \theta^{tT}}{\|\theta^t\|^2} \right) \frac{\partial \mathcal{L}}{\partial \theta}\end{aligned}$$

- Thus, $\mathbf{w} \perp \Delta \mathbf{w}$
- Let $c = \|\Delta \mathbf{w}\| / \|\mathbf{w}\|$, notice: $\|\mathbf{w}^{t+1}\| = \sqrt{(1 + c^2)} \|\mathbf{w}^t\|$
- Large $c \implies$ large $\|\mathbf{w}^{t+1}\| \implies$ small gradient $\Delta \mathbf{w}$
- Small $c \implies \|\mathbf{w}^{t+1}\| \approx \|\mathbf{w}^t\| \implies$ gradient scale $\Delta \mathbf{w}$ unchanged

Effect of backpropagating through normalization:

- Let $\theta = \frac{\gamma(\mathbf{w})}{\|\mathbf{w}\|_2}$, then using SGD:

$$\begin{aligned}\Delta \mathbf{w} &= -\eta \frac{\partial \mathcal{L}}{\partial \mathbf{w}} \\ &= -\eta \frac{\gamma}{\|\mathbf{w}^t\|_2} \left(\mathbf{I} - \frac{\theta^t \theta^{tT}}{\|\theta^t\|^2} \right) \frac{\partial \mathcal{L}}{\partial \theta}\end{aligned}$$

- Thus, $\mathbf{w} \perp \Delta \mathbf{w}$
- Let $c = \|\Delta \mathbf{w}\| / \|\mathbf{w}\|$, notice: $\|\mathbf{w}^{t+1}\| = \sqrt{(1 + c^2)} \|\mathbf{w}^t\|$
- Large $c \implies$ large $\|\mathbf{w}^{t+1}\| \implies$ small gradient $\Delta \mathbf{w}$
- Small $c \implies \|\mathbf{w}^{t+1}\| \approx \|\mathbf{w}^t\| \implies$ gradient scale $\Delta \mathbf{w}$ unchanged

The above analysis applies to NormProp as well!




Figure 1: Training error for CIFAR-10 using different network parameterizations. For *weight normalization*, *batch normalization*, and *mean-only batch normalization* we show results using Adam with a learning rate of 0.003. For the normal parameterization we instead use 0.0003 which works best in this case. For the last 100 epochs the learning rate is linearly decayed to zero.

Model	Test Error
Maxout [6]	11.68%
Network in Network [17]	10.41%
Deeply Supervised [16]	9.6%
ConvPool-CNN-C [26]	9.31%
ALL-CNN-C [26]	9.08%
our CNN, mean-only B.N.	8.52%
our CNN, weight norm.	8.46%
our CNN, normal param.	8.43%
our CNN, batch norm.	8.05%
ours, W.N. + mean-only B.N.	7.31%

Figure 2: Classification results on CIFAR-10 without data augmentation.

Figures taken from Salimans, Tim, and Diederik P. Kingma. "Weight normalization: A simple reparameterization to accelerate training of deep neural networks." Advances in Neural Information Processing Systems. 2016.

Table of Contents

- 1 Introduction
- 2 Motivation
- 3 Batch Normalization
- 4 Normalization Propagation
- 5 Weight Normalization
- 6 Layer Normalization**
- 7 Conclusion

Key idea:

- Compute statistics for each sample separately

Key idea:

- Compute statistics for each sample separately
- Trick: Compute mean and standard deviation across units instead of mini-batch

Key idea:

- Compute statistics for each sample separately
- Trick: Compute mean and standard deviation across units instead of mini-batch
- Aimed towards application to recurrent neural nets:
 - RNNs have variable number of temporal layers
 - There can be more number of layers at test time (BN is not applicable)

- Consider a temporal hidden layer's pre-activation $\mathbf{a}^t = \mathbf{W}_{hh}\mathbf{h}^{t-1} + \mathbf{W}_{xh}\mathbf{x}^t$
- Then layer normalized pre-activation \mathbf{a}^t is given by:

$$\text{LN}(\mathbf{a}^t) = \frac{\gamma}{\sigma^t} \odot (\mathbf{a}^t - \mu^t) + \beta$$

$$\mu^t = \frac{1}{H} \sum_{i=1}^H a_i^t \quad \sigma^t = \sqrt{\frac{1}{H} \sum_{i=1}^H (a_i^t - \mu^t)^2}$$

$$\text{LN}(\mathbf{a}^t) = \frac{\gamma}{\sigma^t} \odot (\mathbf{a}_t - \mu^t) + \beta$$

$$\mu^t = \frac{1}{H} \sum_{i=1}^H a_i^t \quad \sigma^t = \sqrt{\frac{1}{H} \sum_{i=1}^H (a_i^t - \mu^t)^2}$$

Effects of layer normalization:

- Invariance to weight scaling and translation
- Data rescaling and translation
- Norm of weight controls effective learning rate




Figure 2: Validation curves for the attentive reader model. BN results are taken from [Cooijmans et al., 2016].

Figures taken from Ba, Jimmy Lei, Jamie Ryan Kiros, and Geoffrey E. Hinton. "Layer normalization." arXiv preprint arXiv:1607.06450 (2016).

Table of Contents

- 1 Introduction
- 2 Motivation
- 3 Batch Normalization
- 4 Normalization Propagation
- 5 Weight Normalization
- 6 Layer Normalization
- 7 Conclusion**

- Removing internal covariate shift in DNNs leads to faster convergence
- Generally, even good initialization plays an important role
(Mishkin and Matas, 2016)
- Making representations (scale, shift) invariant seems to boost convergence
- SGD + better optimization = better generalization?
– what about bad local minima?

(Zhang et al, 2016)