

# ASSIGNMENT 1: IMPLEMENTATION OF MULTI LAYER PERCEPTRONS [IFT6135]

JOSEPH D. VIVIANO, JONATHAN GUYMONT, MARZIEH MEHDIZADEH

## 1. MLPs ON MNIST

The initial model build contained the following parameters for the layers:

$h_0 = 784$  units (input layer).

$h_1 = 512$  units (hidden layer 1).

$h_2 = 512$  units (hidden layer 2).

$h_3 = 10$  units (output layer, softmax).

This resulted in a model with 669,706 parameters (adding up all weight matrices and biases):

$$n = (512 \times 784) + (512) + (512 \times 512) + (512) + (10 \times 512) + (10)$$

Please see Appendix A for the implementation of the MLP.

**1.1. Initialization Schemes.** We compared training performance of the above MLP under three weight initialization conditions: 1) where all weights are initialized to zeros, 2) where each weight was sampled from a Normal distribution, and 3) where the initial weights were sampled from a uniform distribution using the Glorot initialization scheme.

The above model was trained using ReLU nonlinearities, a learning rate of 0.01, and a mini-batch size of 100 (shuffled), for 10 epochs. We used Stochastic Gradient Descent (SGD) optimizer and cross entropy as our training criterion. Please see the Appendix B for the implementation of the training loop.

**Figure 1** shows that the three initialization methods produced extremely different behaviour in the loss function over epochs.

Initialization using all zeros produced no learning, as the weights were unable to break symmetry. Initialization using normally distributed weights produced very large losses during the first few epochs before producing good results. This is likely because the weight initialization used too many large values, producing large gradients. Glorot initialization produced the best (lowest) losses immediately and at the end of training. This is because the Glorot initialization scheme aims to stabilize the activation variances and the back-propagated gradient variance, which leads to

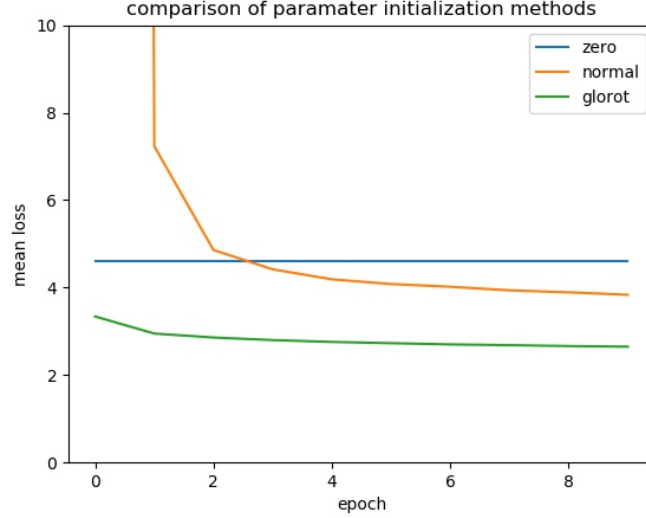


FIGURE 1. Comparison of initialization schemes.

faster convergence.

**1.2. Learning Curves of Models of Differing Capacity.** In this section, we compared the learning curves of two models of differing capacity. Both models were trained with a batch size of 100 (shuffled), learning rate of 0.1, Glorot weight initialization, ReLU nonlinearities, and SGD with cross entropy loss over 100 epochs. The small model had 266610 parameters:

$h_0 = 784$  units (input layer).  
 $h_1 = 300$  units (hidden layer 1).  
 $h_2 = 100$  units (hidden layer 2).  
 $h_3 = 10$  units (output layer, softmax).

**figure 2** shows the training and validation curves for this model:

The large model had 545810 parameters:

$h_0 = 784$  units (input layer).  
 $h_1 = 500$  units (hidden layer 1).  
 $h_2 = 300$  units (hidden layer 2).  
 $h_3 = 10$  units (output layer, softmax).

**figure 3** shows the training and validation curves for this model:

As can be seen, the larger model confers no benefit on the validation set even though the training set accuracy is better. This is because the model with larger




FIGURE 2. Training curve of the small model.

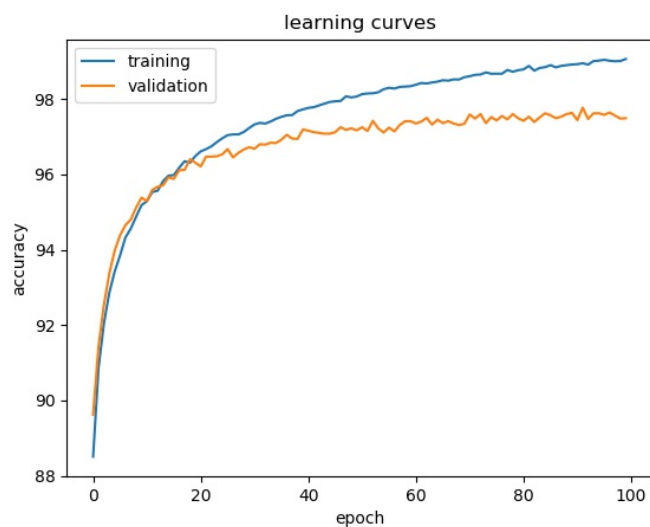


FIGURE 3. Training curve of the large model.

capacity tends to fit the random noise in the training data, which provides no information about the true relationship between the inputs and the targets. Therefore, there is no improvement of the generalization error. When the capacity of the model is high enough to cause the algorithms to memorize random noise in the training set, we say that the model has high variance because different training data are

likely to produce a different model.

**1.3. Training Set Size, Generalization Gap, and Standard Error.** In this section, we compared the effect of training set size on the generalization gap (difference between validation set accuracy and test set accuracy). We used the large model from the previous section. All models were trained with a batch size of 100 (shuffled), learning rate of 0.1, Glorot weight initialization, ReLU nonlinearities, and SGD with cross entropy loss over 100 epochs.

We ran 5 tests, calculating the generalization gap when training with 1%, 2%, 5%, 10%, or 100% of the training data. For each test, we ran 5 training sessions using the above parameters. The results of these tests are shown below:

trainpct	trial1	trial2	trial3	trial4	trial5	mean	se
1	15.54	15.15	16.06	14.4	15.31	15.292	0.2708763556
2	9.98	11.8	10.58	10.73	10.18	10.654	0.3165059241
5	6.59	6.65	6.27	7.11	6.88	6.7	0.1414213562
10	4.77	4.73	5.02	5.07	4.83	4.884	0.0680881781
100	1.604	1.444	1.59	1.344	1.48	1.4924	0.048217839

As can be seen, with larger training set sizes, the generalization gap narrows dramatically (from approximately 15% difference in accuracy, down to a nearly 1% difference in accuracy). Furthermore, the standard error of the mean across the 5 trials also decreases, showing less variability in outcomes across the 5 training sessions when larger training set sizes are used. We conclude that larger training set sizes produce more stable and accurate models with respect to their generalization performance (i.e., test set performance). From a bias-variance point of view, since the mean and the standard deviation of the generalization gap are at their lowest when the training set size is the highest, we could also say that increasing the size of the training set decreases both bias and variance (assuming the added training set data truly comes from the same distribution as the test and validation set data.).

## 2. MLPs ON 20 NEWSGROUPS

**2.1. Preprocessing.** In this section we use the similar MLP as described before, but this time with only one hidden layer of 100 units, and SGD used a momentum of 0.9, Glorot initialization was used in all cases, as well as a batch size of 20. We compared tf-idf and z-scoring preprocessing for text data with raw data, by analyzing their training and test accuracies on the data using all three methods (**figure 4**).

First we observe that tf-idf preprocessing affords a learning rate one order of magnitude greater than either using the raw data, or z-score normalized data. We speculate that this is because the loss when analyzing tf-idf data is informed more by high-information words. Both raw and z-score normalized text data give high weight to common words like ‘and’, ‘or’, and ‘to’, which carry no information about document type since they are common across the entire corpus. Therefore the error gradient of models that include these low-information words in prediction should be much steeper, as these low-information words would lead to a lot of very bad predictions. In contrast, tf-idf downweights the contribution of common words,




FIGURE 4. Training curves for three preprocessing strategies.

and emphasizes regularly-appearing words in the document, but not in the corpus, thereby focusing on features with high information relevant to the task.

Second, if  $\epsilon$  was zero during z-scoring, we could end up with NaNs in our data set if the standard deviation of any word count was zero (i.e., that word always occurs in documents the exact same number of times). This could be fixed by replacing NaNs with some constant value, including zero. Another way to minimize this problem would be to use something like min-max normalization (i.e, scale all counts between

0-1), which is guaranteed not to divide your counts by zero in any reasonable case.

Tf-idf’s advantage, as explored above, is that it down-weights common words that occur in all document types, allowing the model to focus on words unique to each document type. This will result in a better document classifier.

**2.2. Variance in Training.** Figure 5 shows the loss for the first 5000 minibatch updates for a model trained on tf-idf preprocessed data, with a learning rate of 0.2 and a minibatch size of either 1 or 100:

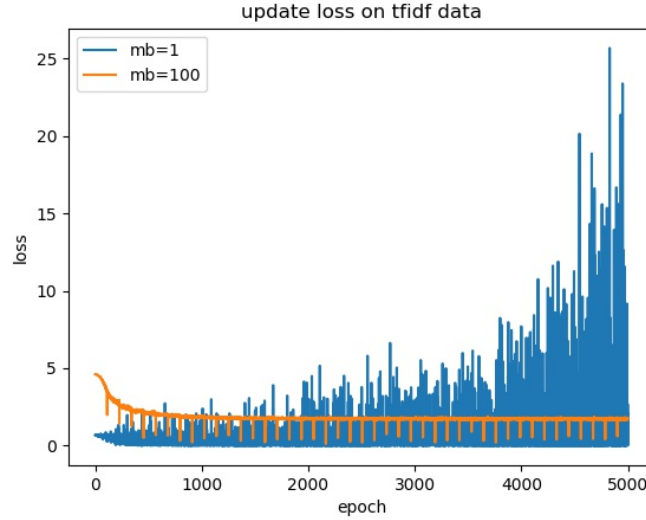


FIGURE 5. Minibatch size and the loss.

While both models have updates of the same magnitude, training performance is very different. First, the gradient computed on single training examples (in the minibatch=1 case) are likely to vary greatly from one another, so each update is going to push the model in different directions. This results in the high variance observed in this case. In contrast, the results seen in the minibatch=100 case are effectively the average loss across 100 minibatch=1 cases, which explains the lower variance. Furthermore, this plot is misleading. Since the minibatch=100 loss plotted is the result of  $100 \times$  the training data of the minibatch=1 loss, the end loss is much lower for the minibatch=100 case because that case has trained on  $100 \times$  the training data.

We could reduce the variance of the minibatch=1 case simply by taking smaller steps each update, that is, reducing the learning rate. Since each loss computed on a single training sample is a noisy estimate of the population loss given the model, the gradient is likely to be inaccurate. Therefore, large steps taken dependent on these gradients are likely to be less productive than more incremental steps.

## APPENDIX A. MLP MODEL

The MLP was implemented as a pytorch `Module` class. It accepts a list of hidden layer sizes, allowing the user to specify an MLP of any length using 3 variables: `h0` = input size, `hn` = hidden layer sizes (as a list), and `ho` = output size.

This class also contains an `initializer` method for initializing weights and `count_params` method for counting all model parameters).

---

```

Class MnistMLP(torch.nn.Module):
    """MLP classifier for MNIST"""
    def __init__(self, h0, hn, ho):
        """
        h0 — input size (flat)
        hn — a list of all hidden layer sizes
        ho — output size (number of classes; flat)
        """
        super(MnistMLP, self).__init__()

        # input —> hid1
        architecture = [torch.nn.Linear(h0, hn[0]), torch.nn.ReLU()]

        # hidden layers
        for i in range(1, len(hn)):
            architecture.append(torch.nn.Linear(hn[i-1], hn[i]))
            architecture.append(torch.nn.ReLU())

        # output
        architecture.append(torch.nn.Linear(hn[-1], ho))

        # use nn to define model
        self.mlp = torch.nn.Sequential(*architecture)
        self.clf = torch.nn.LogSoftmax(dim=0)

    def forward(self, X):
        return(self.clf(self.mlp(X).squeeze()))

    def initializer(self, init_type='glorot'):
        """
        model — a pytorch sequential model
        init_type — one of 'zero', 'normal', 'glorot'

        Takes in a model, initializes it to all-zero, normal distribution
        sampled, or glorot initialization. Glorot == xavier.
        """
        if init_type not in ['zero', 'normal', 'glorot']:
            raise Exception('init_type invalid')

        for k, v in self.mlp.named_parameters():
            if k.endswith('weight'):
                if init_type == 'zero':
                    torch.nn.init.constant(v, 0)
                elif init_type == 'normal':
                    torch.nn.init.normal(v)
                elif init_type == 'glorot':
                    torch.nn.init.xavier_uniform(v, gain=calculate_gain('relu'))

```

```

        else:
            raise Exception('invalid_init_type')

def count_params(self):
    """
    Returns a count of all parameters
    """
    param_count = 0
    for k, v in self.mlp.named_parameters():
        param_count += np.prod(np.array(v.size()))

    return(param_count)

```

---

## APPENDIX B. TRAINING LOOP

The implementation of the training loop can be seen below, and is used throughout the report:

---

```

CUDA = torch.cuda.is_available()

def run_experiment(clf, lr, epochs, loaders, momentum=False, init_type='
    glorot'):

    if len(loaders) == 3:
        train, valid, test = loaders
    elif len(loaders) == 2:
        train, valid = loaders
        test = None
    else:
        raise Exception('loaders_malformed')

    clf.initializer(init_type=init_type)

    if CUDA:
        clf = clf.cuda()

    if momentum:
        optimizer = torch.optim.SGD(clf.parameters(), lr=lr, momentum=0.9)
    else:
        optimizer = torch.optim.SGD(clf.parameters(), lr=lr)

    #lossfn = torch.nn.CrossEntropyLoss() # don't use! B/C I specify
    #LogSoftmax
    lossfn = torch.nn.NLLLoss()

    epoch_loss, valid_acc, train_acc = [], [], []
    best_valid_acc, gen_gap = 0, 0

    all_losses = []
    for ep in range(epochs):

        epoch_losses = []

        # training data
        for batch_idx, (X_train, y_train) in enumerate(train):

```



```

    if CUDA:
        X_train, y_train = X_train.cuda(), y_train.cuda()

    # initialize batch
    optimizer.zero_grad()
    X_train, y_train = Variable(X_train), Variable(y_train)

    # make predictions — flatten each image (batchsize x pixels)
    train_pred = clf.forward(X_train.view(X_train.shape[0], -1))

    # calculate loss (cross entropy)
    loss = lossfn(train_pred, y_train)
    epoch_losses.append(loss.data[0])
    all_losses.append(loss.data[0])

    # calculate dloss/dx for all parameters that have
    # requires_grad
    loss.backward()

    # update parameter values
    optimizer.step()

    # average loss for epoch
    epoch_loss.append(np.mean(epoch_losses))

    # validation accuracy for this epoch
    this_valid_acc = evaluate(clf, valid)
    valid_acc.append(this_valid_acc)

    # training accuracy for this epoch
    this_train_acc = evaluate(clf, train)
    train_acc.append(this_train_acc)

    # keep track of the best validation accuracy, generalization gap
    if this_valid_acc > best_valid_acc:
        best_valid_acc = this_valid_acc

    if test:
        this_test_acc = evaluate(clf, test)
        gen_gap = this_train_acc - this_test_acc

    # update every n epochs
    if (ep+1) % 5 == 0:
        curr_lr = optimizer.state_dict()['param_groups'][0]['lr']
        print(' + [{:03d}] loss={:0.6f} acc={:0.2f} / {:0.2f} lr={:0.5f}'
              .format(
                  ep+1, epoch_loss[-1], train_acc[-1], valid_acc[-1],
                  curr_lr))

results = {'clf': clf,
          'epoch_loss': epoch_loss,
          'all_loss': all_losses,
          'train_acc': train_acc,
          'valid_acc': valid_acc,
          'best_valid_acc': best_valid_acc,
          'gen_gap': gen_gap}
    
```

```
return(results)
```

---

UNIVERSITÉ DE MONTRÉAL

*Email address:* joseph@viviano.ca, jonathan.guymont@com, marzieh.mehdizadeh@gmail.com