

ASSIGNMENT 1: THEORY OF MULTI LAYER PERCEPTRONS

[IFT6135]

JOSEPH D. VIVIANO

1. NEURAL NETWORK CLASSIFICATION

- (a) $\sigma(z) = 1/(1 + e^{-z})$ is an appropriate activation function for binary classification (i.e., the sigmoid activation function).
- (b) $\sigma(z)$ (above) outputs numbers on the interval $[0, 1]$ and can be interpreted as a probability.
- (c) $L = -y \log g(z) - (1 - y) \log(1 - g(z))$ is the cross entropy loss.
- (d) For cross entropy loss, $df/da = g(a)(1 - g(a))$, and $\partial L/\partial a = f - y$
- (e) $L = 1/2(y - f)^2$ is the mean squared error loss (assuming a single example).
- (f) For mean squared error loss, $\partial L/\partial a = -(y - f)g(a)(1 - g(a))$
- (g) For binary classification, we want our responses to be bounded between two values, the extremes representing each class and the distance to each representing our certainty of choosing one class versus the other class. The cross entropy loss outputs a probability, which fulfills the aforementioned requirements, while the mean squared error loss is unbounded, making it a less suitable loss for binary classification. One could define a threshold on which to binarize the outputs to produce a classification, but if we were to use a sigmoid output unit (or another option with similar properties), the units would saturate quickly and very small gradients (and therefore little learning) would be produced if the neuron's mean squared error loss ever grew too large.

2. NEURAL NETWORK REPRESENTATION

Fill the parameters of a MLP binary classifier using a Heavyside step function, two input neurons, three hidden neurons, and one output neuron.

Three decision boundaries are required to separate all samples by their class. First we define four relevant points:

$$\begin{aligned} p_1 &= (-4, 0) \\ p_2 &= (0, -4) \end{aligned}$$

Date: Feburary 2018.

$$\begin{aligned} p_3 &= (4, 0) \\ p_4 &= (0, 4) \end{aligned}$$

And second, we define the three decision boundaries:

$$\begin{aligned} d_1 &= \overrightarrow{p_1 p_2} \\ d_2 &= \overrightarrow{p_2 p_3} \\ d_3 &= \overrightarrow{p_3 p_4} \end{aligned}$$

These decision boundaries are computed by the hidden layer of our network. Each neuron i receives two inputs: x_{1i} and x_{2i} modulated by weights w_{1i} and w_{2i} . d_1 has slope = -1 and passes through $y = -4$, d_2 has slope = 1 and passes through $y = -4$, and d_3 has slope = -1 and passes through $y = 4$. Therefore the three decision planes can be expressed using the form $y = mx + b$ as so:

$$\begin{aligned} d_1 &= -1x - 4 \\ d_2 &= 1x - 4 \\ d_3 &= -1x + 4 \end{aligned}$$

Since the weights define the slope of each decision boundary, they are trivial to derive:

$$\begin{aligned} w_{11} &= -1 \text{ and } w_{21} = 1 \text{ for } d_1, \\ w_{12} &= 1 \text{ and } w_{22} = 1 \text{ for } d_2, \\ w_{13} &= -1 \text{ and } w_{23} = 1 \text{ for } d_3. \end{aligned}$$

The bias of each decision boundary defines its offset from the origin, or the euclidean distance between the origin and the closet point of the decision boundary. Therefore, the biases for these decision boundaries b_1, b_2, b_3 are set to be -2, -2, and 2, respectively.

Finally, these decision boundaries must be added and subtracted from one another using the output layer weights u_1, u_2, u_3 , and the bias of the output neuron c set's the default state of the decision space. Therefore, to carve out the required boundary, we set the output biases to $u_1 = -1$ to remove the lower boundary, $u_2 = -1$ to remove the side boundary, and $u_3 = 1$ to add the top boundary, which leaves us with the correct decision boundary. The output bias u must be set $-1 < u < 0$ so that these subtractions and additions produce the correct behaviour after the Heavyside step function, which sets all values less than 0 to 0, and all values greater than 0 to 1.

Please see **figure 1** for an illustration incorporating these parameters.

3. ACTIVATION FUNCTIONS

(a) Show that $ReLU'(x)$ is given by the Heavyside step function $H(x)$.

The heavyside step function is always 0 when $x < 0$, always 1 when $x > 0$, and vertical (undefined) at $x = 0$. $ReLU(x) = \max(0, x)$ is 0 when $x < 0$, some linear

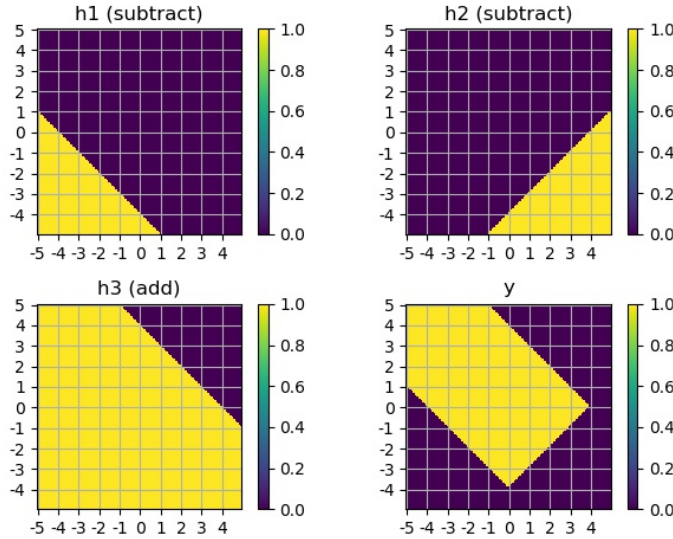


FIGURE 1. Heavyside decision boundary.

function $y = f(x)$ when $x > 0$, and 0 when $x = 0$. Therefore the derivative of the ReLU function is defined everywhere except at $x = 0$, since that one location is not continuous. Furthermore when $x < 0$ $dy/dx = 0$ since it is constant, and when $x > 0$, $dy/dx = 1$ since it is the derivative of a single variable. Assembling these pieces, we get the heavyside step function.

(b) Two definitions of ReLU using the Heavyside step function include $xH(x)$ and $x^{H(x)} - \mathbb{1}_{x \leq 0}$.

(c) Show the sigmoid function behaves asymptotically like $H(x)$.

As $\lim_{x \rightarrow \infty} 1/(1+e^{-x})$, $\sigma(x)$ approaches 1 because the term e^{-x} quickly approaches 0, and conversly, as $\lim_{x \rightarrow -\infty} 1/(1+e^{-x})$, $\sigma(x)$ approaches 0 because the denominator grows very large. Therefore $\sigma(x)$ behaves asymptotically as the $H(x)$ function.

(d) Show $H'(x)$ is given by the dirac delta function.

First, observe that the Heavyside function is constant everywhere, except at $x = 0$. Therefore $H'(x < 0) = 0$ and $H'(x > 0) = 0$. Second, at $x = 0$, the slope of the Heavyside function is ∞ (a vertical line), so $H'(x = 0) = \infty$. If we combine these three intervals, $x < 0$, $x = 0$, and $x > 0$, we get the dirac function, which is equal to $H'(x)$:

$$(3.1) \quad H'(x) = \delta(x) = \begin{cases} +\infty, & x = 0 \\ 0, & x \neq 0 \end{cases}.$$

If we integrate over any area where the entire density of the distribution is concentrated at $x=0$, we end up with the second property of the Dirac function since the sum of the area under an integral is always equal to 1:

$$(3.2) \quad \int_{-\infty}^{\infty} H'(x) dx = \int_{-\infty}^{\infty} \delta(x) dx = 1.$$

4. GRADIENTS AND NETWORKS

(a) For the softmax function $\sigma(\mathbf{z})_j = \frac{e^{z_j}}{\sum_{k=1}^K e^{z_k}}$, the jacobian is $J = (\partial \text{softmax} / \partial z)_{ij}$.

$$(4.1) \quad \frac{\partial \text{softmax}_i(z)}{\partial z_j} = \frac{\delta_{ij} e^{z_i} \sum_k e^{z_k} - e^{z_i} e^{z_j}}{(\sum_k e^{z_k})^2}$$

Which simplifies to:

$$(4.2) \quad \frac{\partial \text{softmax}_i(z)}{\partial z_j} = \text{softmax}_i(z)(\delta_{ij} - \text{softmax}_j(z)).$$

In both cases δ_{ij} is the Kronecker delta, which evaluates to 1 if $i = j$ and is 0 otherwise.

(b)

Observe the jacobian as a matrix equation. For brevity we represent *softmax* as *s* below:

$$\begin{bmatrix} s_1(1 - s_1(z)) & -s_1 s_2 & \cdots & -s_1 s_n \\ -s_2 s_1 & s_2(1 - s_2(z)) & \cdots & -s_2 s_n \\ \vdots & \vdots & \ddots & \vdots \\ s_n(1 - s_n(z)) & -s_n s_2 & \cdots & s_n(1 - s_n(z)) \end{bmatrix}$$

(c) The jacobian of the logistic sigmoid applied element-wise to a vector x is the diagonal of the above matrix, i.e., $\text{diag}(\text{softmax}_i(z)(\delta_{ij} - \text{softmax}_j(z)))$.

(d) To see that $J_y(x)g_y = g_y \circ \sigma'(x)$, first expand g_y to the matrix $J_y(g_{y1}, \dots, g_{yn})^T$, and multiply $\sigma'(x)$ into each element: $(\sigma'(x_1)g_{y1}, \dots, \sigma'(x_n)g_{yn})^T$. From this, we see $(g_{y1}, \dots, g_{yn})^T \odot (\sigma'(x_1), \dots, \sigma'(x_n))^T$.

If $J_y(x)$ were the jacobian of the softmax function, this does not hold because J_y would not be a diagonal matrix.

5. SOFTMAX ACTIVATION FUNCTION

(a) The softmax is invariant under translation, i.e., when a constant c is added to all values of x , i.e., $\text{softmax}(x + c) = e^{x_i + c} / \sum_k e^{x_k + c}$, because it is normalized by the sum of all e^x , i.e., if all values are equally shifted by some constant c , they will

simple be range-normalized back to the range $[0, 1]$.

(b) The softmax function is not invariant under multiplication, because when $-1 < x < 1$, multiplication will shrink $|x|$, whereas when $x > 1 \vee x < -1$, multiplication will increase $|x|$. Multiplication would have the effect of making the softmax more certain of the most likely output, i.e., it would sharpen the distribution of probabilities computed by the softmax.

(c) The sigmoid function is equivalent to a 2-class softmax. Let me show you:

First, observe the sigmoid: $\sigma(z) = 1/(1 + e^{-z})$. Isn't it nice? Second, observe the general form of the k-class softmax: $\frac{e^{x_i}}{\sum_{i=1}^k e^{x_i}}$.

Now let's split up the softmax into the two-class case:

$$(5.1) \quad \text{softmax}(x) = \left(\frac{e^{x_1}}{e^{x_1} + e^{x_2}} + \frac{e^{x_2}}{e^{x_1} + e^{x_2}} \right)$$

So for the x_1 case:

$$(5.2) \quad \text{softmax}(x_1) = \frac{e^{x_1}}{e^{x_1} + e^{x_2}} = \frac{1}{1 + e^{x_2 - x_1}} = \frac{1}{1 + e^{-x_1 - x_2}} = \frac{1}{1 + e^{-z}}$$

Therefore if $z = x_1 - x_2$, $\sigma(z) = \text{softmax}(x_1)$. We can compute a similar equality for $\text{softmax}(x_2)$:

$$(5.3) \quad \text{softmax}(x_2) = \frac{e^{x_2}}{e^{x_1} + e^{x_2}} = \frac{1}{e^{x_1 - x_2} + 1} = 1 - \frac{1}{1 + e^{-x_1 - x_2}} = 1 - \frac{1}{1 + e^{-z}}$$

Therefore $1 - \sigma(z) = \text{softmax}(x_2)$, and more completely:

$$(5.4) \quad (\sigma(z), 1 - \sigma(z)) = (\text{softmax}(x_1), \text{softmax}(x_2))$$

(d) Note that in the previous 2-class example, we divide each element of the softmax vector by its numerator, which results in the term $1 + e^{-x_1 - x_2}$ in the denominator. It is not hard to see that, for the three class case, if we were to divide the fraction by the same numerator, we would be left with $1 + e^{-x_1 - x_2} + e^{-x_1 - x_3}$ in the denominator. We could therefore assign two variables $a = -x_1 - x_2, b = -x_1 - x_3$, to represent the 3 class softmax as 2 variables a and b . By induction, we can see that the k -class softmax can be represented as using variables $a, b, \dots, k - 1$.

6. USING CROSS ENTROPY (NEGATIVE LOG LIKELIHOOD) FOR REAL-VALUED DATA

(a) *Derive the cross-entropy cost function using the maximum likelihood principle for x*

We set $\text{prob}(x = 1) = p$, and $\text{prob}(x = 0) = 1 - p$

This can be expressed as the bernoulli distribution $f(x; p) : p^x(1-p)^{1-x}$ for $x \in 0, 1$.

This allows us to express this as the negative log likelyhood seen in question 1: $-x \log p + (1 - x) \log(1 - p)$, aka the cross entropy loss.

(b) Here I present a probabilistic interpretation of cross-entropy using the KL divergence between two distributions. According to wikipedia, the relationship between Cross-entropy and K-L divergence is:

$$(6.1) \quad H(p, q) = \mathbb{E}_p[-\log q] = H(p) + D_{\text{KL}}(p||q) = -\sum_x p(x) \log q(x)$$

Where p and q are two probability distributions (meaning $p \in [0, 1], q \in [0, 1]$). p represents the likelihood of x truly being 1, and q represents the probability of predicting that x is 1. $H(p)$ is the entropy of p , and $H(p, q)$ is the cross-entropy of these two probability distributions. Therefore, the cross entropy is simply the sum of the $H(p)$, and the K-L divergence $p||q$, a measure of the difference between the distribution of model predictions q and the true distribution p .

The K-L divergence for these two distributions is:

$$(6.2) \quad D_{\text{KL}}(p||q) = -\sum_x p(x) \log \frac{q(x)}{p(x)}$$

and the entropy for p is:

$$(6.3) \quad H(p) = \sum_x p(x) \log p(x)$$

therefore it is trivial to show:

$$(6.4) \quad \sum_x p(x) \log p(x) - \sum_x p(x) \log \frac{q(x)}{p(x)} = -\sum_x p(x) \log q(x)$$

7. Q7: DERIVING THE GLOROT INITIALIZATION SCHEME

The Glorot initialization scheme was developed to work on the sigmoid activation function, assuming all weights were in the linear regime of the sigmoid. The goal of this initialization scheme is to ensure that all layers have a $\mathbb{E} = 0$ and $\text{Var}(X) = 1$.

The values in layer Y will be the result of all n layers before it:

$$(7.1) \quad Y_i = W_1X_1 + W_2X_2 + \cdots + W_nX_n$$

Where X_i is the inputs to that layer and W_i are the weights into that layer. The variance of W_iX_i :

$$(7.2) \quad \text{Var}(W_iX_i) = E[X_i]^2\text{Var}(W_i) + E[W_i]^2\text{Var}(X_i) + \text{Var}(W_i)\text{Var}(X_i)$$

Since we want the mean to be zero, we can remove all terms with E:

$$(7.3) \quad \text{Var}(W_iX_i) = \text{Var}(W_i)\text{Var}(X_i)$$

Next we assume that these are uncorrelated variables, so the variance of Y is simply the sum of all n input variances:

$$(7.4) \quad \text{Var}(Y_i) = \sum_n \text{Var}(W_i)\text{Var}(X_i)$$

If we want $\text{Var}(Y_i)$ to be equal to $\text{Var}(X_i)$, we therefore need $\text{Var}(W_i)$ to be equal to $\frac{1}{n}$, so the terms cancel out.

This is true for both the forward pass and backward pass of backpropagation independently, with each serving as an input to layer Y , we simply take the sum of the inverse of the weight counts for the forward and backward pass (n_i and n_{i+1}):

$$(7.5) \quad \text{Var}(W_i) = \frac{2}{n_i + n_{i+1}}$$

In the case of ReLU, the linear assumption used in the original paper does not hold (recall the paper only analyzed the linear regime of the sigmoid activation function around zero). With ReLU, all values below $x = 0$ are 0, while all values above are x , which effectively cuts our variance in half since our initialized values are centered around zero. Therefore we must double our weight variance to keep the variance constant throughout the layers:

$$(7.6) \quad \text{Var}(W_i) = \frac{4}{n_i + n_{i+1}}$$

UNIVERSITÉ DE MONTRÉAL

Email address: joseph@viviano.ca